



操作系统实验

Linker & Loaders
(Based on practice)

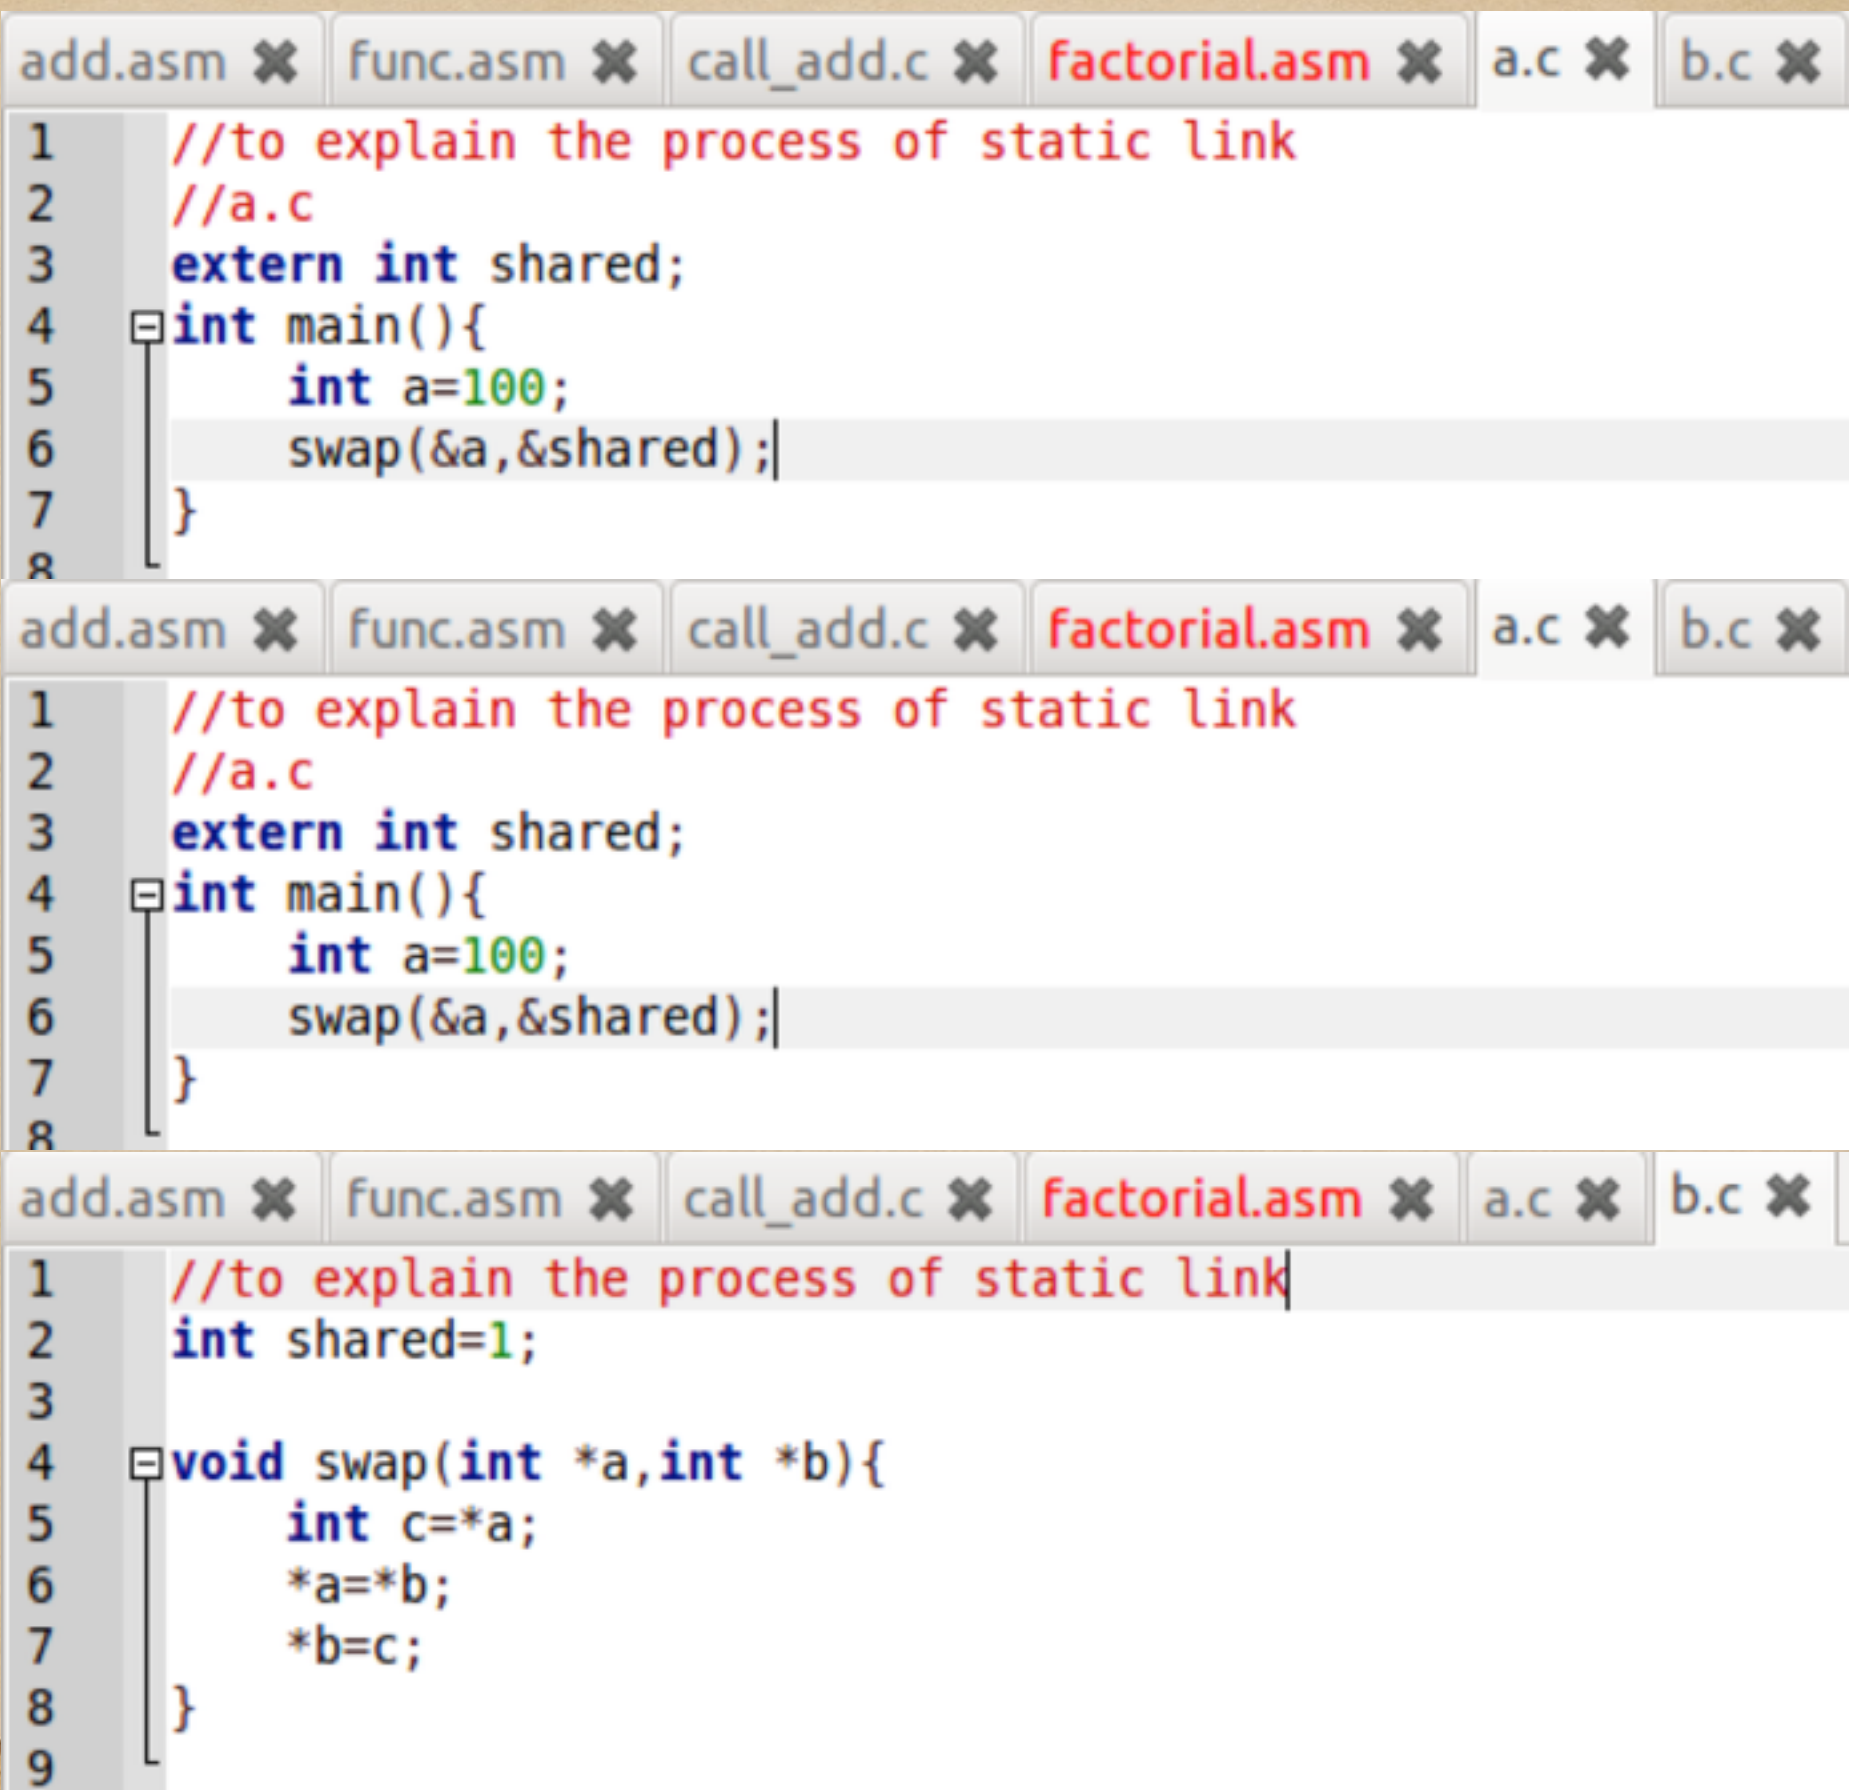
刘璟 ljing12@software.nju.edu.cn

2015.4.17

Outline

- ◆ static linking
- ◆ dynamic linking
- ◆ Something about loader

Static linking



```
add.asm x func.asm x call_add.c x factorial.asm x a.c x b.c x
1 //to explain the process of static link
2 //a.c
3 extern int shared;
4 int main(){
5     int a=100;
6     swap(&a,&shared);
7 }
8

add.asm x func.asm x call_add.c x factorial.asm x a.c x b.c x
1 //to explain the process of static link
2 //a.c
3 extern int shared;
4 int main(){
5     int a=100;
6     swap(&a,&shared);
7 }
8

add.asm x func.asm x call_add.c x factorial.asm x a.c x b.c x
1 //to explain the process of static link
2 int shared=1;
3
4 void swap(int *a,int *b){
5     int c=*a;
6     *a=*b;
7     *b=c;
8 }
9
```


Static linking

```
$ objdump -h a.o
```

```
a.o:      file format elf32-i386
```

Sections:

Idx	Name	Size	VMA	LMA	File off	Algn
0	.text	00000027	00000000	00000000	00000034	2**2
	CONTENTS, ALLOC, LOAD, RELOC, READONLY, CODE					
1	.data	00000000	00000000	00000000	0000005c	2**2
	CONTENTS, ALLOC, LOAD, DATA					
2	.bss	00000000	00000000	00000000	0000005c	2**2
	ALLOC					
3	.comment	0000002b	00000000	00000000	0000005c	2**0
	CONTENTS, READONLY					
4	.note.GNU-stack	00000000	00000000	00000000	00000087	2**0
	CONTENTS, READONLY					
5	.eh_frame	00000038	00000000	00000000	00000088	2**2
	CONTENTS, ALLOC, LOAD, RELOC, READONLY, DATA					

Static linking

```
$ objdump -h b.o
```

```
b.o:      file format elf32-i386
```

Sections:

Idx	Name	Size	VMA	LMA	File off	Algn
0	.text	00000022	00000000	00000000	00000034	2**2
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
1	.data	00000004	00000000	00000000	00000058	2**2
	CONTENTS, ALLOC, LOAD, DATA					
2	.bss	00000000	00000000	00000000	0000005c	2**2
	ALLOC					
3	.comment	0000002b	00000000	00000000	0000005c	2**0
	CONTENTS, READONLY					
4	.note.GNU-stack	00000000	00000000	00000000	00000087	2**0
	CONTENTS, READONLY					
5	.eh_frame	00000038	00000000	00000000	00000088	2**2
	CONTENTS, ALLOC, LOAD, RELOC, READONLY, DATA					

Static linking

```
$ ld a.o b.o -e main -o ab
```

```
$ objdump -h ab
```

```
ab:      file format elf32-i386
```

Sections:

Idx	Name	Size	VMA	LMA	File off	Algn
0	.text	0000004a	08048094	08048094	00000094	2**2
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
1	.eh_frame	00000058	080480e0	080480e0	000000e0	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
2	.data	00000004	08049138	08049138	00000138	2**2
	CONTENTS, ALLOC, LOAD, DATA					
3	.comment	0000002a	00000000	00000000	0000013c	2**0
	CONTENTS, READONLY					

What happened?

The similar segments merge in the output file.

Static linking

\$objdump -d a.o

```
a.o:      file format elf32-i386
```

```
Disassembly of section .text:
```

```
00000000 <main>:
```

0:	55	push	%ebp
1:	89 e5	mov	%esp,%ebp
3:	83 e4 f0	and	\$0xfffffffff0,%esp
6:	83 ec 20	sub	\$0x20,%esp
9:	c7 44 24 1c 64 00 00	movl	\$0x64,0x1c(%esp)
10:	00		
11:	c7 44 24 04 00 00 00	movl	\$0x0,0x4(%esp)
18:	00		
19:	8d 44 24 1c	lea	0x1c(%esp),%eax
1d:	89 04 24	mov	%eax,(%esp)
20:	e8 fc ff ff ff	call	21 <main+0x21>
25:	c9	leave	
26:	c3	ret	

What happened in the two boxes?

Static linking

\$objdump -d ab

```
08048094 <main>:
8048094:      55                push    %ebp
8048095:      89 e5             mov     %esp,%ebp
8048097:      83 e4 f0          and     $0xfffffffff0,%esp
804809a:      83 ec 20          sub     $0x20,%esp
804809d:      c7 44 24 1c 64 00 00 movl    $0x64,0x1c(%esp)
80480a4:      00
80480a5:      c7 44 24 04 38 91 04 movl    $0x8049138,0x4(%esp)
80480ac:      08
80480ad:      8d 44 24 1c       lea     0x1c(%esp),%eax
80480b1:      89 04 24          mov     %eax,(%esp)
80480b4:      e8 03 00 00 00    call   80480bc <swap>
80480b9:      c9               leave
80480ba:      c3               ret
80480bb:      90               nop

080480bc <swap>:
80480bc:      55                push    %ebp
80480bd:      89 e5             mov     %esp,%ebp
80480bf:      83 ec 10          sub     $0x10,%esp
80480c2:      8b 45 08          mov     0x8(%ebp),%eax
80480c5:      8b 00             mov     (%eax),%eax
80480c7:      89 45 fc          mov     %eax,-0x4(%ebp)
80480ca:      8b 45 0c          mov     0xc(%ebp),%eax
80480cd:      8b 10             mov     (%eax),%edx
80480cf:      8b 45 08          mov     0x8(%ebp),%eax
80480d2:      89 10             mov     %edx,(%eax)
80480d4:      8b 45 0c          mov     0xc(%ebp),%eax
80480d7:      8b 55 fc          mov     -0x4(%ebp),%edx
80480da:      89 10             mov     %edx,(%eax)
80480dc:      c9               leave
80480dd:      c3               ret
```

relocation

Static linking

How to relocate? Relocation Table & Symbol Table

`$objdump -r a.o`

```
a.o:          file format elf32-i386

RELOCATION RECORDS FOR [.text]:
OFFSET      TYPE          VALUE
00000015  R_386_32              shared
00000021  R_386_PC32            swap

RELOCATION RECORDS FOR [.eh_frame]:
OFFSET      TYPE          VALUE
00000020  R_386_PC32            .text
```

`$objdump -s a.o`

Symbol table '.symtab' contains 11 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000000	0	FILE	LOCAL	DEFAULT	ABS	a.c
2:	00000000	0	SECTION	LOCAL	DEFAULT	1	
3:	00000000	0	SECTION	LOCAL	DEFAULT	3	
4:	00000000	0	SECTION	LOCAL	DEFAULT	4	
5:	00000000	0	SECTION	LOCAL	DEFAULT	6	
6:	00000000	0	SECTION	LOCAL	DEFAULT	7	
7:	00000000	0	SECTION	LOCAL	DEFAULT	5	
8:	00000000	39	FUNC	GLOBAL	DEFAULT	1	main
9:	00000000	0	NOTYPE	GLOBAL	DEFAULT	UND	shared
10:	00000000	0	NOTYPE	GLOBAL	DEFAULT	UND	swap

Dynamic linking:

The advantage of dynamic linking:

- ◆ Dynamically linked shared libraries are easier to create than static linked shared libraries.
- ◆ Dynamically linked shared libraries are easier to update than static linked shared libraries.
- ◆ The semantics of dynamically linked shared libraries can be much closer to those of unshared libraries.
- ◆ Dynamic linking permits a program to load and unload routines at runtime, a facility that can otherwise be very difficult to provide.

Dynamic linking:

Steps:

- ◆ Starting the dynamic linker
 - ◆ bootstrap
- ◆ Finding the libraries
- ◆ Shared library initialization
 - ◆ relocation & initialization

Dynamic linking:

```
add.asm x func.asm x call_add.c x factorial.asm x a.c x b.c x pa.c x
1 //explain dynamic linking
2 //pa.c
3 #include "Lib.h"
4
5 int main(){
6     func(1);
7     return 0;
8 }
```

```
add.asm x func.asm x call_add.c x factorial.asm x a.c x b.c x pa.c x pb.c
1 //explain the dynamic linking
2 //pb.c
3
4 int main(){
5     func(2);
6     return 0;
7 }
```

```
add.asm x func.asm x call_add.c x factorial.asm x a.c x b.c x pa.c x pb.c
1 //explain dynamic linking
2 //Lib.h
3 #ifndef LIB_H
4 #define LIB_H
5 void func(int i);
6 #endif
```

```
add.asm x func.asm x call_add.c x factorial.asm x a.c x b.c x pa.c x pb.c
1 //explain the dynamic linking
2 //Lib.c
3
4 #include <stdio.h>
5 void func(int i){
6     printf("printing from Lib.so%d\n",i);
7     sleep(-1);
8 }
9
```


Dynamic linking:

```
$ gcc -fPIC -shared -o Lib.so lib.c
```

```
$ gcc -o pa pa.c ./Lib.so
```

```
$ gcc -o pb pb.c ./Lib.so
```

What if run pa and pb simultaneously?

Lib.so will only be loaded to memory once

```
$ ./pa
```

```
$ ./pb
```

```
$ ps -e
```


Dynamic linking:

```
~$ cat /proc/5846/maps
08048000-08049000 r-xp 00000000 08:01 134551 [redacted] /workspace
/asm/prac1/pa
08049000-0804a000 r--p 00000000 08:01 134551 [redacted] /workspace
/asm/prac1/pa
0804a000-0804b000 rw-p 00001000 08:01 134551 [redacted] /workspace
/asm/prac1/pa
b7579000-b757b000 rw-p 00000000 00:00 0
b757b000-b771f000 r-xp 00000000 08:01 1055162 /lib/i386-linux-gnu/libc-2.15.s
o
b771f000-b7721000 r--p 001a4000 08:01 1055162 /lib/i386-linux-gnu/libc-2.15.s
o
b7721000-b7722000 rw-p 001a6000 08:01 1055162 /lib/i386-linux-gnu/libc-2.15.s
o
b7722000-b7725000 rw-p 00000000 00:00 0
b7734000-b7735000 rw-p 00000000 00:00 0
b7735000-b7736000 r-xp 00000000 08:01 134548 [redacted] /workspace
/asm/prac1/Lib.so
b7736000-b7737000 r--p 00000000 08:01 134548 [redacted] /workspace
/asm/prac1/Lib.so
b7737000-b7738000 rw-p 00001000 08:01 134548 [redacted] /workspace
/asm/prac1/Lib.so
b7738000-b773a000 rw-p 00000000 00:00 0
b773a000-b773b000 r-xp 00000000 00:00 0 [vdso]
b773b000-b775b000 r-xp 00000000 08:01 1055152 /lib/i386-linux-gnu/ld-2.15.so
b775b000-b775c000 r--p 0001f000 08:01 1055152 /lib/i386-linux-gnu/ld-2.15.so
b775c000-b775d000 rw-p 00020000 08:01 1055152 /lib/i386-linux-gnu/ld-2.15.so
bf9ad000-bf9ce000 rw-p 00000000 00:00 0 [stack]
```


Dynamic linking:

How to use your codes and experimental results to understand the process of dynamic linking?

- ◆ Waiting for you!
- ◆ Bonus!
- ◆ Good luck!

Hints: use command 'readelf', 'objdump', 'ps' and 'cat' properly.

Thanks~ $O(n_n)O!$